

PROJECT 5 REPORT

PRANIT SEHGAL

ASU ID : 1225456193

REPORT : Comparative Analysis of Convolutional Neural Network Designs for CIFAR-10 Object Classification

Google Colab Link :

<https://colab.research.google.com/drive/1gyIGlzNE2OwDk6WY4ru7b513i9hZs7te?usp=sharing>

A. System Design

Motivation: The rapid advancements in deep learning have led to substantial improvements in image classification tasks. This project arose from a desire to examine how different Convolutional Neural Network (CNN) architectures perform against each other in a controlled setting. The ultimate aim was to gain a deeper understanding of model designs and their implications on performance metrics.

High-level Design: The project journey traversed four distinct phases, each built upon a specific CNN philosophy:

1. **Basic CNN Model:** This foundational phase used a basic design comprising several convolutional layers, serving as a reference point for the following iterations.
2. **Pooling Enhanced Model:** To manage computational constraints and capture intrinsic patterns, pooling layers were integrated into the design.
3. **Dropout Augmented Model:** Recognizing the detrimental effects of overfitting, dropout layers were introduced to encourage model robustness.
4. **MobileNetV2 Implementation:** Representing the pinnacle of lightweight and efficient architectures, MobileNetV2 was the chosen representative of pre-designed CNN architectures.

Observations and Difficulties: As the experimentation progressed, it became evident that even straightforward models can achieve remarkable performance if well-tuned. However, challenges like overfitting and model selection criteria necessitated rigorous iteration and refinement.

B. Experiment and Algorithm

Experimental Framework: Experiments were diligently executed using TensorFlow, a widely acclaimed deep learning platform. The chosen dataset, segmented into training, validation, and test sets, showcased colored images spanning multiple object categories.


Configurations and Justifications:


1. **Phase 1:** A straightforward model was essential to understand the baseline performance without the influence of advanced mechanisms. This phase leveraged three convolutional layers of increasing filter complexity, providing a rudimentary yet effective design.
2. **Phase 2:** Pooling layers were methodically positioned post convolutional layers to extract dominant features and reduce computational demands, thereby aiding in model efficiency.
3. **Phase 3:** Introducing dropout served a dual purpose - preventing neurons from becoming overly specialized and enhancing generalization.
4. **Phase 4:** With its reputation for compactness without compromising performance, MobileNetV2's implementation was to ascertain how pre-designed architectures fare against customized designs.

Hyperparameter choices, predominantly the learning rate and batch size, were governed by best practices, and 'adam' was the optimizer of choice due to its adaptive learning rate capabilities.



C. Results


PHASE 1 :








PHASE 2:








PHASE 3 :



PHASE 4:



Accuracies Overview: The experiments conducted across the four phases yielded the following accuracies:

- Phase 1: 50.55%
- Phase 2: 71.61%
- Phase 3: 70.40%
- Phase 4: 60.68%

Analysis of Results:

1. Phase 1

- Architecture: This phase utilized a basic Convolutional Neural Network (CNN) design without pooling layers.
- Result: The accuracy achieved was 50.55%. Without pooling layers, the network relies solely on the convolutional layers to extract features. This might not be optimal for capturing abstract features, which might explain the modest accuracy.

2. Phase 2

- Architecture: Introduced pooling layers after the convolutional layers.
- Result: There was a notable increase in accuracy to 71.61%. Pooling layers are designed to reduce the spatial dimensions of the data, making the network focus on the most crucial elements. This addition clearly bolstered the model's capacity to generalize from the training data, reflecting in the accuracy jump.

3. Phase 3

- Architecture: Dropouts were integrated into the model to combat potential overfitting.
- Result: The accuracy was slightly reduced to 70.40%. The introduction of dropout layers aims to prevent overfitting by randomly setting a fraction of input units to 0 at each update during training. However, it's possible that the exact dropout rate used might not have been optimal, slightly affecting the performance.

4. Phase 4

- Architecture: Leveraged the ShuffleNet architecture, focusing on training the fully connected layers.
- Result: The accuracy dropped to 60.68%. While ShuffleNet is designed to be efficient for mobile and embedded devices, its architecture might not be directly transferable to the CIFAR-10 dataset without more extensive fine-tuning. Moreover, by freezing the main layers and training only the fully connected ones, we might not be exploiting the full potential of the architecture.

Conclusion: The Phase 2 model with pooling layers showed the best performance among all. While subsequent phases introduced techniques and architectures that have their strengths, for the specific dataset and problem at hand, the CNN model with pooling layers demonstrated superior capability. Future iterations should look into more detailed tuning of the hyperparameters, dropout rates, and perhaps explore even deeper architectures while being wary of overfitting.

Performance Metrics: Performance was gauged across three fronts: training, validation, and test datasets. The evolution of accuracy and loss with epochs painted a vivid picture of model behavior.

Comparison and Visualization: Phase 2 emerged as the star, achieving a commendable validation accuracy of 71.61%. The basic model from Phase 1 secured 50.55%, highlighting that simplicity can occasionally rival complexity. Phase 3, despite its anti-overfitting measures, landed at 70.40%. Surprisingly, MobileNetV2, despite its prowess, clocked 60.68%.

Graphical representations, including epoch-wise accuracy/loss plots and confusion matrices, elucidated finer model performance nuances, shedding light on strengths and weaknesses.

D. Discussions

Overview: Phase 2's design did best in our comparison, but all the different designs had their unique features. It was interesting to see how MobileNetV2 performed, making us question if famous models always work best everywhere.

Challenges with Classifying Objects: It's tough to classify objects because there are so many different types, some look very similar, and there's always the risk of the model memorizing the training data (overfitting).

What's Next?: In the future, we could try new ways to prevent overfitting, add more variety to our training data, or combine different models for better results. We might also fine-tune our models more carefully, possibly using advanced tuning techniques.

Which Objects Were Hard to Recognize?: Our data showed that some objects, like cats and dogs, got mixed up more often. This hints at areas where we could focus our research and maybe develop special techniques for certain object types.

Concluding Reflections

This project taught us a lot, showing both what we know about deep learning and what we still have to discover. It reminded us how important it is to keep trying different approaches, test our models thoroughly, and stay updated with the latest research in the world of neural networks.